

Oracle Partitioning – part 2

- ◆ Interval Partitioning
 - Create
 - Alter
- ◆ Virtual column based Partitioning
- ◆ Reference Partitioning
- ◆ Partition for Tiered Storage
- ◆ Summary
- ◆ Lab: Partitioning Extensions

Partitioning Extensions

Page 1


Basic Partitioning in Oracle Database

PARTITIONING STRATEGY	DATA DISTRIBUTION	SAMPLE BUSINESS CASE
Range Partitioning	Consecutive ranges of values.	Orders table range partitioned by order_date
List Partitioning	Unordered lists of values.	Orders table list partitioned by country
Hash Partitioning	Internal hash algorithm	Orders table hash partitioned by customer_id
Composite Partitioning <ul style="list-style-type: none">• Range-Range• Range-List• Range-Hash• List-List• List-Range• List-Hash• Hash-Hash• Hash-List• Hash-Range	Combination of two of the above-mentioned basic techniques of Range, List, and Hash	<p>Orders table is range partitioned by order_date and sub-partitioned by hash on customer_id</p> <p>Orders table is range partitioned by order_date and sub-partitioned by range on shipment_date</p> <p>Orders table is list partitioned by country and sub-partitioned by range on order_date</p> <p>Orders table is list partitioned by country and sub-partitioned by hash on customer_id</p>

Page 2

Automate the partition management

- ◆ Partitioning is **key-enabling** functionality for managing large volumes of data
 - One logical object for application transparency
 - Multiple physical segments for administration
- ◆ **Improves Manageability, Availability, and Performance**
- ◆ **BUT** Physical segmentation requires additional data management overhead
 - E.g. new partitions must be created on-time for new data



Page 3

Oracle Database New Partitioning Enhancements

Oracle provides **partitioning extensions** that enhance the usage of the basic partitioning strategies. Partitioning extensions enhance the manageability of partitioned objects and provide more flexibility in defining the partitioning key of a table or even groups of tables that are logically connected through parent-child relationships.

- ◆ **One consistent way to manage all your data**
 - Not just for data warehouse and high-end OLTP any more
 - New referential, virtual column, and interval partitioning features bring partitioning to mainstream
- ◆ **Easier management of today's rapidly growing datasets**
- ◆ **Improved performance**
 - Partition elimination speeds table scans
- ◆ **Reduced costs**
 - Automatically place less used data on lower cost storage

Page 4


□ Interval Partitioning

- ◆ Managing the creation of new partitions can be a cumbersome and highly repetitive task. This is especially true for predictable additions of partitions covering small ranges, such as adding new daily partitions.
- ◆ Interval partitioning automates this operation by creating partitions **on-demand**.
- ◆ This new partitioning feature fully automates the partition creation for **range**.
 - New partitions will be created when they are needed.
 - By defining the interval criteria, the database knows when to create new partitions for new or modified data.
- ◆ You can create single-level interval partitioned tables as well as the following composite partitioned tables:
 - **Interval-range**
 - **Interval-hash**
 - **Interval-list**

Page 5

Interval Partitioning

- ◆ **Automatic interval partitioning**
 - Partitions are created automatically as data arrives
 - Example: Create new partition every month



Page 6

Extension of range partitioning

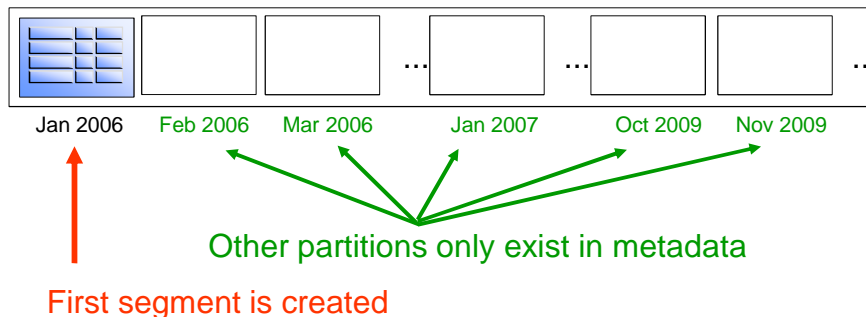
- ◆ Interval partitioning is an **extension of range partitioning** which instructs the database to automatically create partitions of a specified interval when data inserted into the table exceeds all of the existing range partitions.
- ◆ You must specify at least **one range partition**. The range partitioning key value determines the high value of the range partitions, which is called the transition point, and the database creates interval partitions for data beyond that transition point.
- ◆ The lower boundary of every interval partition is the non-inclusive upper boundary of the previous range or interval partition.
- ◆ For example, if you create an interval partitioned table with monthly intervals and the transition point at January 1, 2007, then the lower boundary for the January 2007 interval is January 1, 2007. The lower boundary for the July 2007 interval is July 1, 2007, regardless of whether the June 2007 partition was already created.
- ◆ You can only specify **one partitioning key column**, and it must be of **NUMBER** or **DATE** type.
- ◆ Interval partitioning is not supported for index-organized tables.
- ◆ You cannot create a domain index on an interval-partitioned table.

Page 7

Interval Partitioning


```
CREATE TABLE sales (order_date DATE, ...)
PARTITION BY RANGE (order_date)
INTERVAL(MonthInt(1,'month'))
(PARTITION p_first VALUES LESS THAN ('01-JAN-2006'));
```

Table SALES




Page 8

Interval Partitioning



↑ New segment is automatically allocated

```
INSERT INTO sales (order_date DATE, ...)
VALUES ('04-MAR-2006',...);
```



... whenever data for a new partition arrives ↑

```
INSERT INTO sales (order_date DATE, ...)
VALUES ('17-OCT-2009',...);
```



Page 9

Interval Partitioning

◆ Interval partitioned table can have classical range and automated interval section

- Automated new partition management plus full partition maintenance capabilities: *"Best of both worlds"*

Table SALES



Range partition section → Interval partition section


Transition point

Page 10

Add Interval Partitioning to an existing Range Partition Table

◆ Range partitioned tables can be extended into interval partitioned tables

- Simple metadata command
- Investment protection



```
ALTER TABLE sales (order_date DATE, ...)
SET INTERVAL(MonthInt(1,'month'));
```

Page 11

Interval vs. Range Partitioning

◆ Partition bounds

- Interval partitions have lower and upper bound
- Range partitions only have upper bounds
 - ◆ Lower bound derived by previous partition

◆ Partition naming

- Interval partitions cannot be named in advance
- Range partitions must be named

◆ Partition merge

- Multiple non-existent interval partitions are silently merged
- Only two adjacent range partitions can be merged at any point in time

Page 12

Deferred Segment Creation vs. Interval Partitioning

◆ Interval Partitioning

- Maximum number of one million partitions are pre-defined
 - ◆ Explicitly defined plus interval-based partitions
- No segments are allocated for partitions without data
 - ◆ New record insertion triggers segment creation
- **Ideal for “ever-growing” tables**

◆ “Standard” Partitioning with deferred segment creation

- Only explicitly defined partitions are existent
 - ◆ New partitions have to be added via DDL
- No segments are allocated for partitions without data
 - ◆ New record insertion triggers segment creation when data matches pre-defined partitions
- Ideal for sparsely populated pre-defined tables

Page 13

Interval Partitioning Summary

◆ Interval Partitioning

- Extension to Range Partitioning
- Full automation for equi-sized range partitions
- ◆ **Partitions are created as metadata information only**
 - Start Partition is made persistent
- ◆ **Segments are allocated as soon as new data arrives**
 - No need to create new partitions
 - Local indexes are created and maintained as well
 - No need for any partition management

Page 14

❑ Virtual Column-Based Partitioning

- ◆ In Oracle Database 11g *and later*, you can now partition key columns defined on virtual columns of a table.
- ◆ Frequently, business requirements to logically partition objects does not match existing columns in a one-to-one manner.
- ◆ Oracle partitioning has been enhanced to allow a partitioning strategy being defined on virtual columns, thus enabling a more comprehensive match of the business requirements.
 - Virtual columns are defined by evaluating an expression the results of which become the metadata of the columns for tables.
 - Virtual columns can be defined at table creation or modification time.
 - Virtual columns enable application developers to define computations and transformations as the column (metadata) definition of tables without space consumption.
- ◆ This makes application development easier and less error-prone, as well as enhances query optimization by providing additional statistics to the optimizer for these virtual columns.

Page 15

Virtual Columns

Business Problem

- ◆ Extended Schema attributes are fully derived and dependent on existing common data
- ◆ Redundant storage or extended view definitions are solving this problem today
 - Requires additional maintenance and creates overhead

Solution

- ◆ Oracle Database 11g introduces virtual columns
 - Purely virtual, meta-data only
- ◆ Treated as real columns except no DML
 - Virtual columns can have statistics
 - Virtual columns are eligible as partitioning key
- ◆ Enhanced performance and manageability


Page 16

Virtual Column-Based Partitioning Example

- ◆ REGION requires **no storage**
- ◆ Partition by ORDER_DATE, REGION

ORDERS

ORDER_ID	ORDER_DATE	CUSTOMER_ID...
9834- US -14	12-JAN-2007	65920
8300- EU -97	14-FEB-2007	39654
3886- EU -02	16-JAN-2007	4529
2566- US -94	19-JAN-2007	15327
3699- US -63	02-FEB-2007	18733



REGION AS (SUBSTR(ORDER_ID,6,2))

US

EU

EU

US

US

Page 17

Virtual Columns - Example

- ◆ Base table with all attributes ...
- ... is extended with the virtual (derived) column
- ... and the virtual column is used as partitioning key

```
CREATE TABLE accounts
(acc_no      number(10)    not null,
 acc_name    varchar2(50)  not null,
 ...
```

```
CREATE TABLE accounts
(acc_no      number(10)    not null,
 acc_name    varchar2(50)  not null,
 acc_branch  number(2) generated always as
              (to_number(substr(to_char(acc_no),1,2)))
partition by list (acc_branch)
```

12500	Adams	12
12507	Blake	12
12666	King	12
12875	Smith	12

...

32320	Jones	32
32407	Clark	32
32758	Hura	32
32980	Phillips	32

Page 18

Examples

```
CREATE TABLE emp_year_sal
(ename VARCHAR2(20),
sal NUMBER,
yearly_sal AS (sal*12) VIRTUAL)
PARTITION BY RANGE (yearly_sal)
(PARTITION low_sal VALUES LESS THAN (20000),
PARTITION mid_sal VALUES LESS THAN (40000),
PARTITION high_sal VALUES LESS THAN (60000),
PARTITION others VALUES LESS THAN (MAXVALUE));
```

```
SQL> SELECT ename,sal,yearly_sal FROM emp_year_sal;
```

ENAME	SAL	YEARLY_SAL
SMITH	800	9600
ALLEN	1600	19200
WARD	1250	15000
JONES	2975	35700
MARTIN	1250	15000
BLAKE	2850	34200
CLARK	2450	29400
SCOTT	3000	36000

Page 19

```
SQL> SELECT ename,sal,yearly_sal FROM emp_year_sal PARTITION
(low_sal);
```

ENAME	SAL	YEARLY_SAL
SMITH	800	9600
ALLEN	1600	19200
WARD	1250	15000
MARTIN	1250	15000
TURNER	1500	18000
ADAMS	1100	13200

```
SQL> SELECT ename,sal,yearly_sal FROM emp_year_sal
PARTITION(mid_sal);
```

ENAME	SAL	YEARLY_SAL
JONES	2975	35700
BLAKE	2850	34200
CLARK	2450	29400
SCOTT	3000	36000
FORD	3000	36000

Page 20

REF Partitioning

Business Problem

- ◆ **Related tables benefit from same partitioning strategy**
 - Sample order – lineitem
- ◆ **Redundant storage of the same information solves this problem**
 - Data overhead
 - Maintenance overhead

Solution: REF Partitioning

◆ Intuitive modelling

- Child table inherits the partitioning strategy of parent table through PK-FK relationship
- The partitioning key is resolved through an existing parent-child relationship, enforced by active primary key or foreign key constraints.

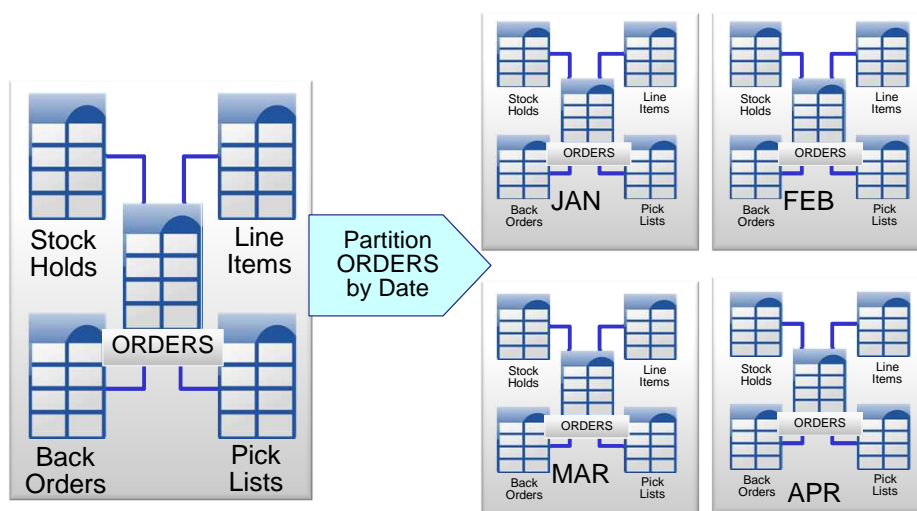
◆ Enhanced Performance and Manageability

- The benefit of this feature is that tables with a parent-child relationship can be logically equi-partitioned by inheriting the partition key from the parent table without duplicating the key columns.
- The logical dependency will also automatically cascade partition maintenance operations, thus making application development easier and less error-prone.

Page 21

REF Partitioning

◆ Inherit partitioning strategy



Page 22

REF Partitioning

◆ Inheritance of partition key couples parent and child tables together

- Child tables do not have a partitioning key
- PK – FK relationship cannot be disabled or even dropped

◆ Due to the tight coupling, some things are different

- [Sub]Partition names are inherited down from the parent to the child tables
 - ◆ No system-generated names unless parent has them
- Child partitions are by default co-located with the parent partition
 - ◆ Default for user is automatically overwritten

Page 23

REF Partitioning

◆ Partition maintenance operations (PMOPs)

- PMOP that change the table structure are implicit for child tables and inherited from the parent table
 - ◆ ADD, DROP, MERGE, and SPLIT
- PMOPs without structure changes are fully supported
 - ◆ MOVE, EXCHANGE
- TRUNCATE works in the presence of PK-FK relationship


◆ Partition-wise Joins (PWJ)

- Joining parent and child tables are always eligible for PWJ, due to the known data co-location in the joining partitions

Page 24

Before REF Partitioning


Table ORDERS



- ◆ RANGE(**order_date**)
- ◆ Primary key **order_id**

- Redundant storage of order_date
- Redundant maintenance

Table LINEITEMS




- ◆ RANGE(**order_date**)
- ◆ Foreign key **order_id**

Page 25

After REF Partitioning

Table ORDERS




- ◆ RANGE(**order_date**)
- ◆ Primary key **order_id**

PARTITION BY REFERENCE

- Partitioning key inherited through PK-FK relationship

Table LINEITEMS



- ◆ RANGE(**order_date**)
- ◆ Foreign key **order_id**

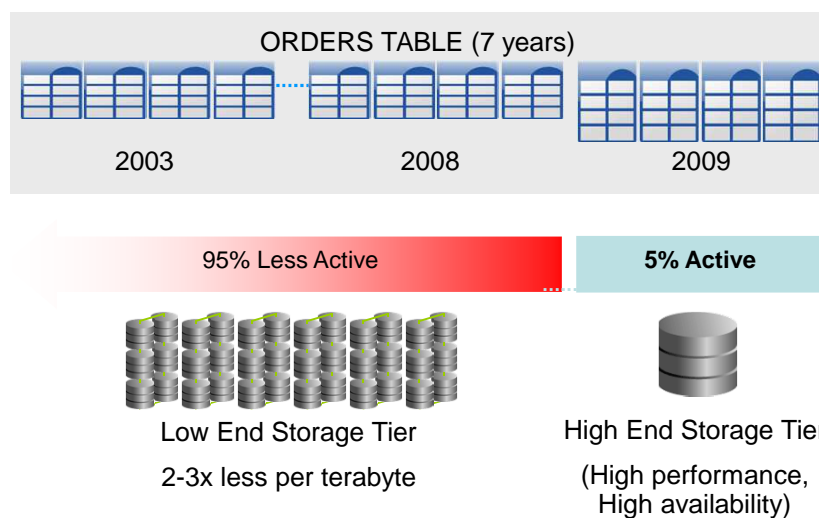
Page 26

New Partitioning Extensions In Oracle 11g/12c

PARTITIONING EXTENSION	DESCRIPTION	SAMPLE BUSINESS CASE
Interval Partitioning <ul style="list-style-type: none"> Interval Interval-Range Interval-List Interval-Hash 	Extension to Range Partition. Defined by an interval, providing equi-width ranges. With the exception of the first partition all partitions are automatically created on-demand when matching data arrives.	Orders table partitioned by order_date with a predefined daily interval, starting with '01-Jan-2013'
Reference Partitioning	Partitioning for a child table is inherited from the parent table through a primary key – foreign key relationship. The partitioning keys are not stored in actual columns in the child table.	(Parent) Orders table range partitioned by order_date and inherits the partitioning technique to (child) order lines table. Column order_date is only present in the parent orders table
Virtual column based Partitioning	Defined by any partition techniques where the partitioning key is based on a virtual column. Virtual columns are not stored on disk and only exist as metadata.	Orders table has a virtual column that derives the sales region based on the first three digits of the customer account number. The orders table is then list partitioned by sales region.

Page 27

Partition for Tiered Storage



Page 28

Summary

◆ Partitioning

- An independent functionality that is beneficial for every environment
- Transparent to the application – no SQL changes
 - ◆ All of the data COULD be queried
- Increases Performance and Scalability on large tables
- Provides facilities for better Manageability and Availability
 - ◆ Tables can be split into many pieces.
 - ◆ Only a subset of the data is queried
 - ◆ Re-orgs & backups can be done on a partition level

◆ When to Use Which Partitioning Method

- First determine if you need to partition the table.
- Next decide which table partitioning method is right for your situation.
- Determine how volatile the data is.
 - ◆ How often are there inserts, updates and deletes?
- Choose your indexing strategy: global or local partitioned indexes.
 - ◆ Each type has its own maintenance consideration.

Page 29